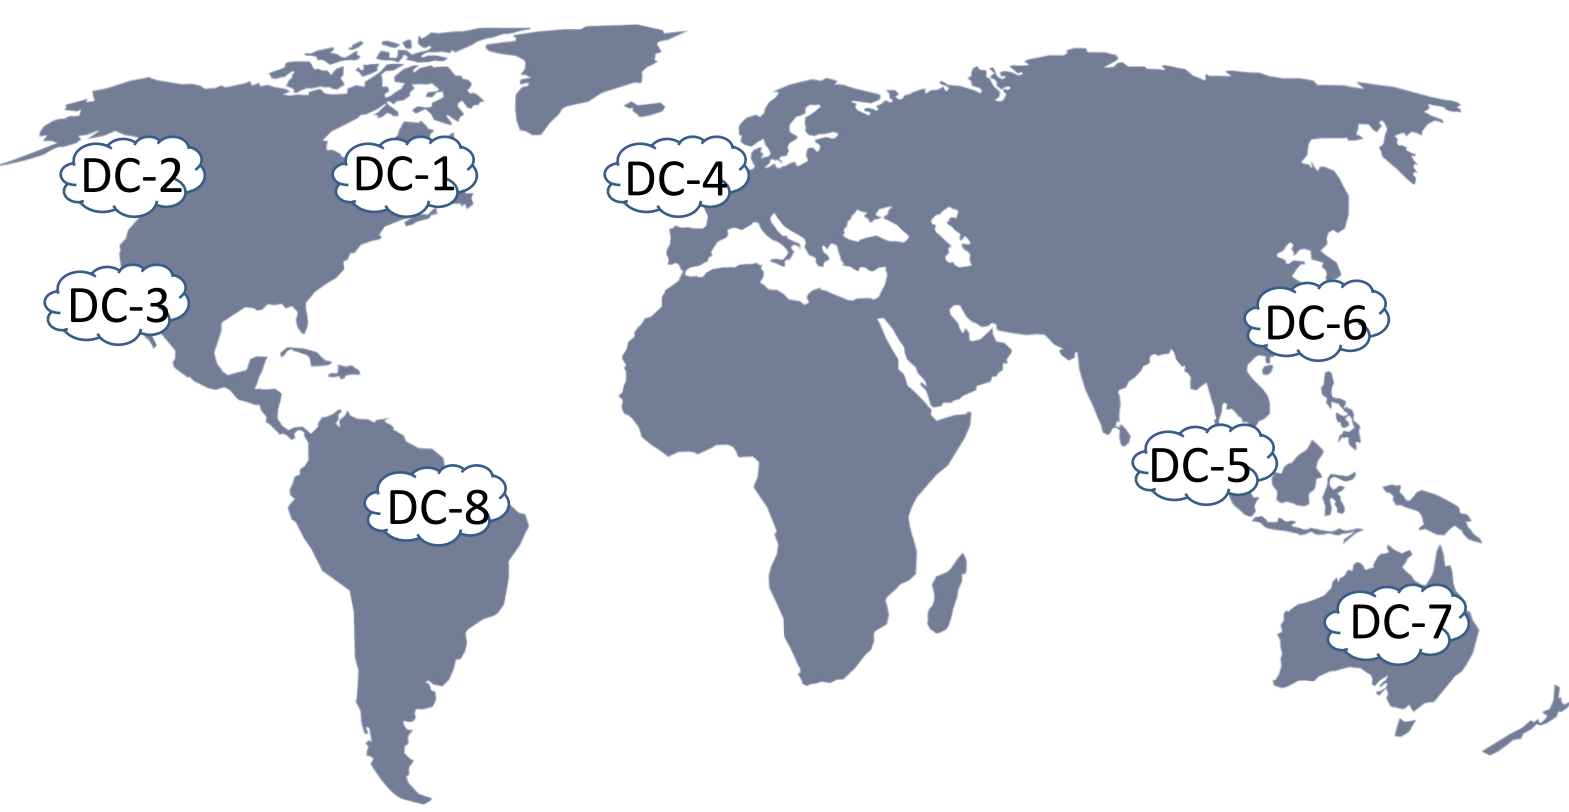


Motivation

Modern Cloud Storage Systems




- Span multiple geo-distributed DCs
 - Twitter, Facebook, Google, Amazon
 - Amazon has 8 worldwide DCs
- Replicate data for low latency
- Use asynchronous write propagation
- Have to handle failures, and partitions


The Consistency Spectrum

Replication, asynchronous write propagation create ordering issues

- Weak “eventually consistent” systems
Widely deployed, but ordering can be confusing
- Strong ordering of all reads and writes across all clients
CAP Theorem \Rightarrow unavailable on partition



Problems with Existing Causal Approaches

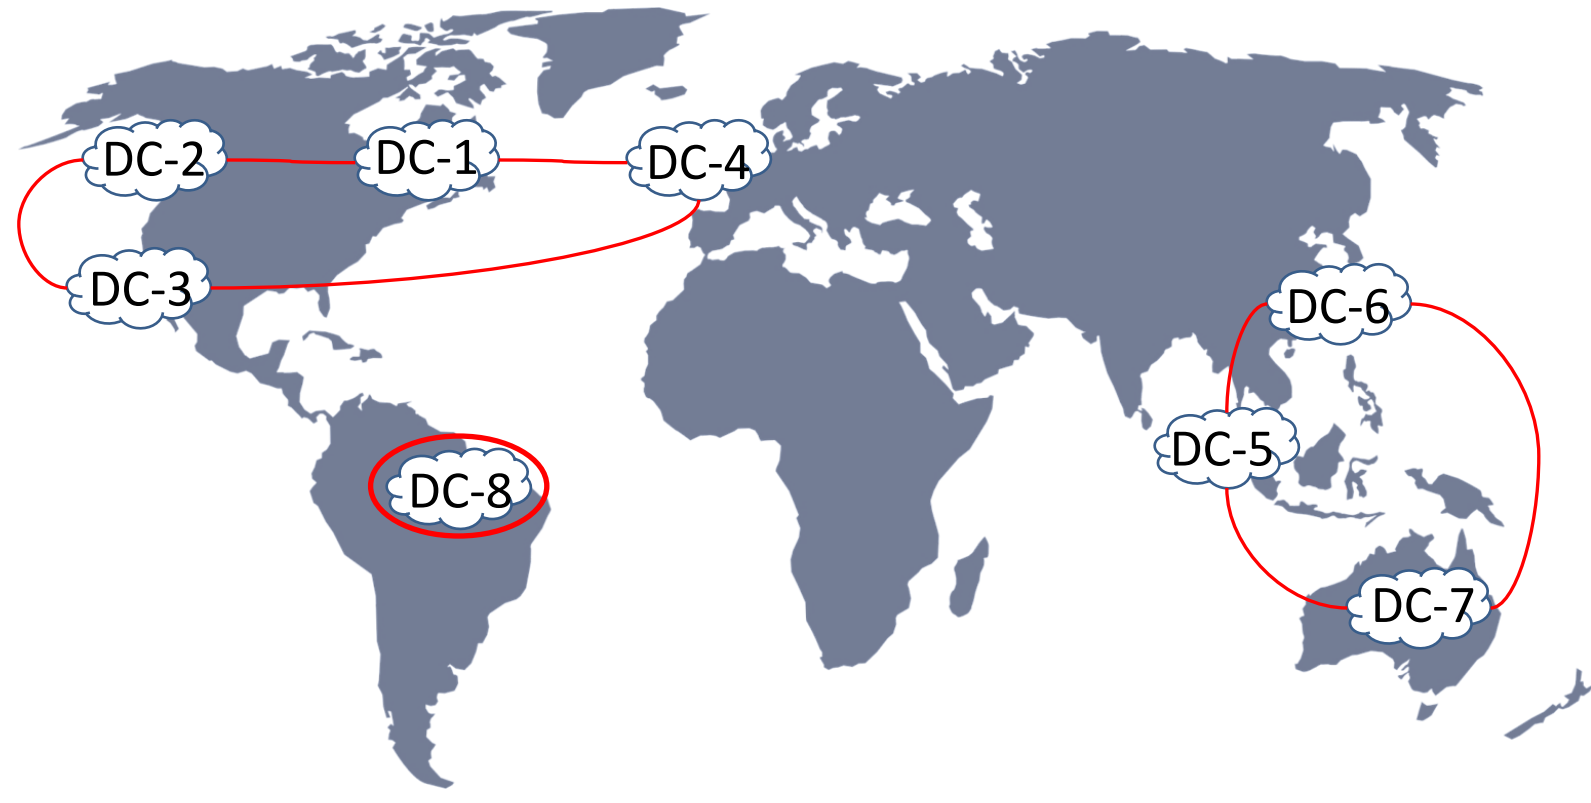


- Static binding:** A user is allowed to access only one DC
- Full replication:** Expensive, scalability issues
- Simple solutions do not work
 - Spreading data across DCs \Rightarrow Availability issues
 - Allowing users to switch DCs \Rightarrow Consistency violation

Karma: First causally consistent geo-replicated cloud storage system with partial replication while preserving consistency and availability

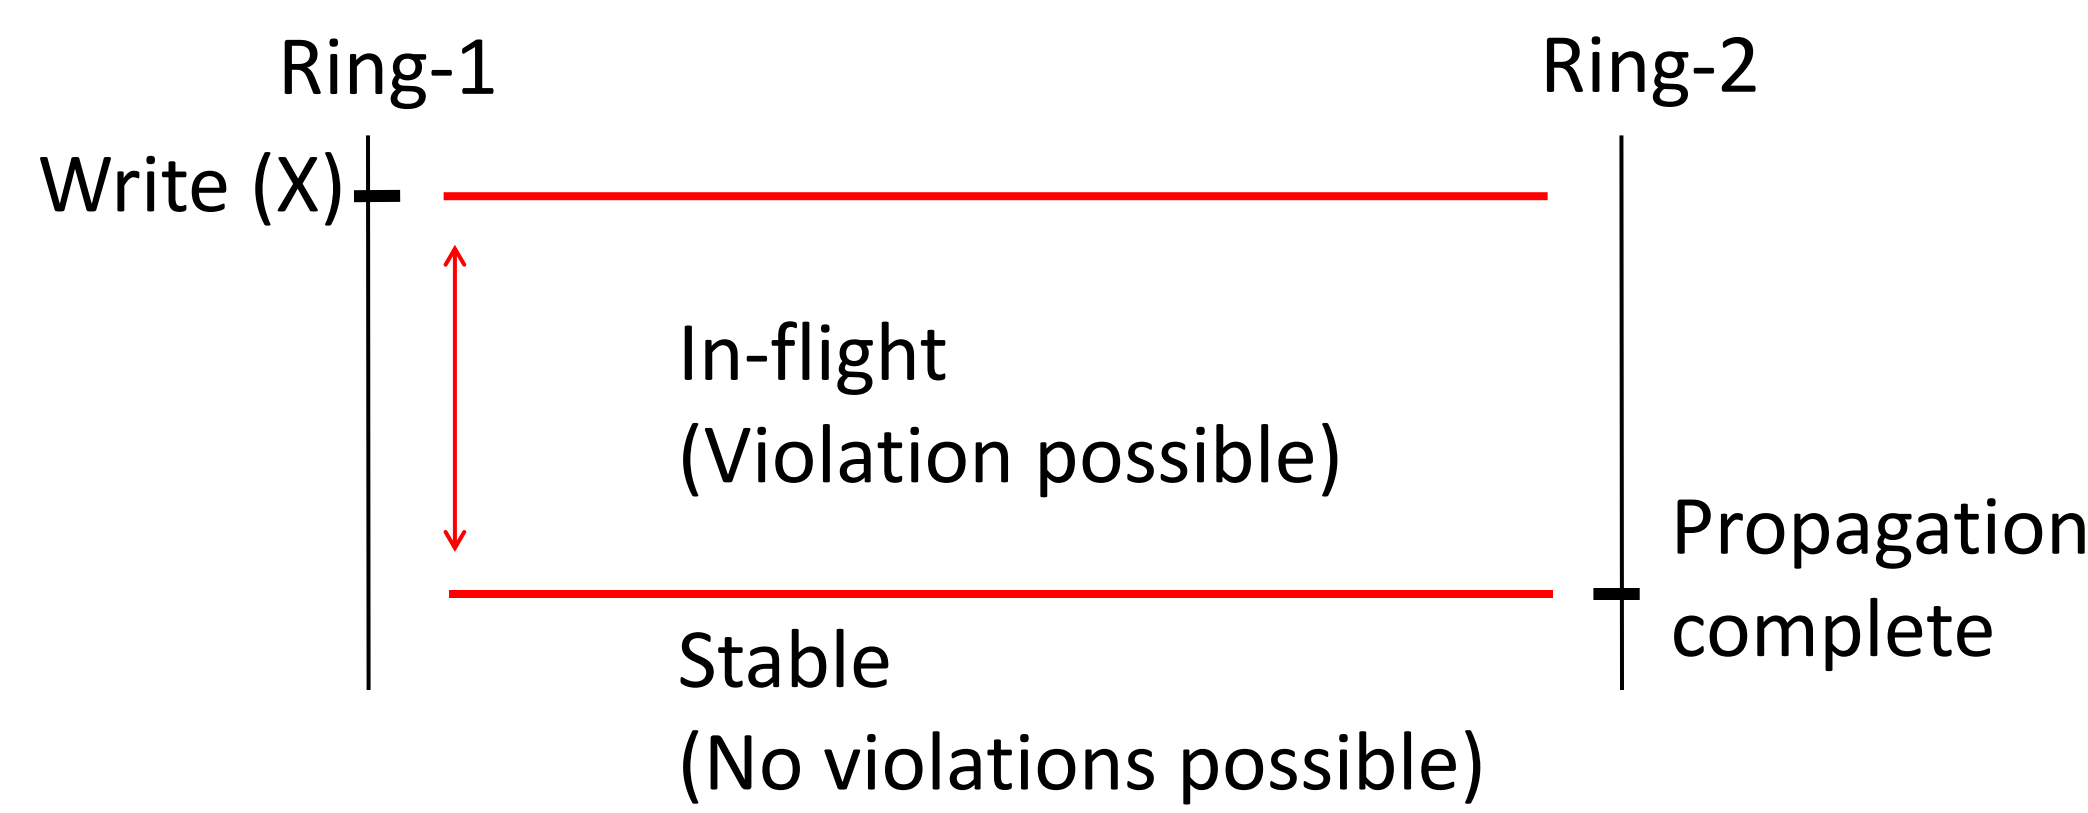
Karma's Key Ideas

Partial Replication



- Decouple rings and DCs
- Rings span multiple DCs
Each ring contains full replica of dataset
- Availability in wide-area rings guaranteed by causality-preserving dynamic ring binding
- DC level caching used for fast reads of remote objects

Dynamic Ring Binding



- Karma's novel mechanism: Dynamic Ring Restrictions (DRR)
- If a client reads an in-flight object from Ring-1
Temporarily restrict client to read all objects from Ring-1
Client can access any ring once in-flight objects are stable

Caching/Write Buffers

- Partial replication \Rightarrow **Remote objects, slow**
- DC-level storage caches enable fast reads:
 - Problem:** Normal cache operation violates causality
 - Solution:** Stable value caching
- Persistent thread-private write buffers enable fast writes
 - All writes are local
 - Reads check write-buffer to avoid violations

Performance Evaluation (R/W : 95/5)

Experimental Setup

- 64-node testbed on PROBE cluster
- 8 data centers, 8 nodes each
- Amazon AWS emulation using DummyNet

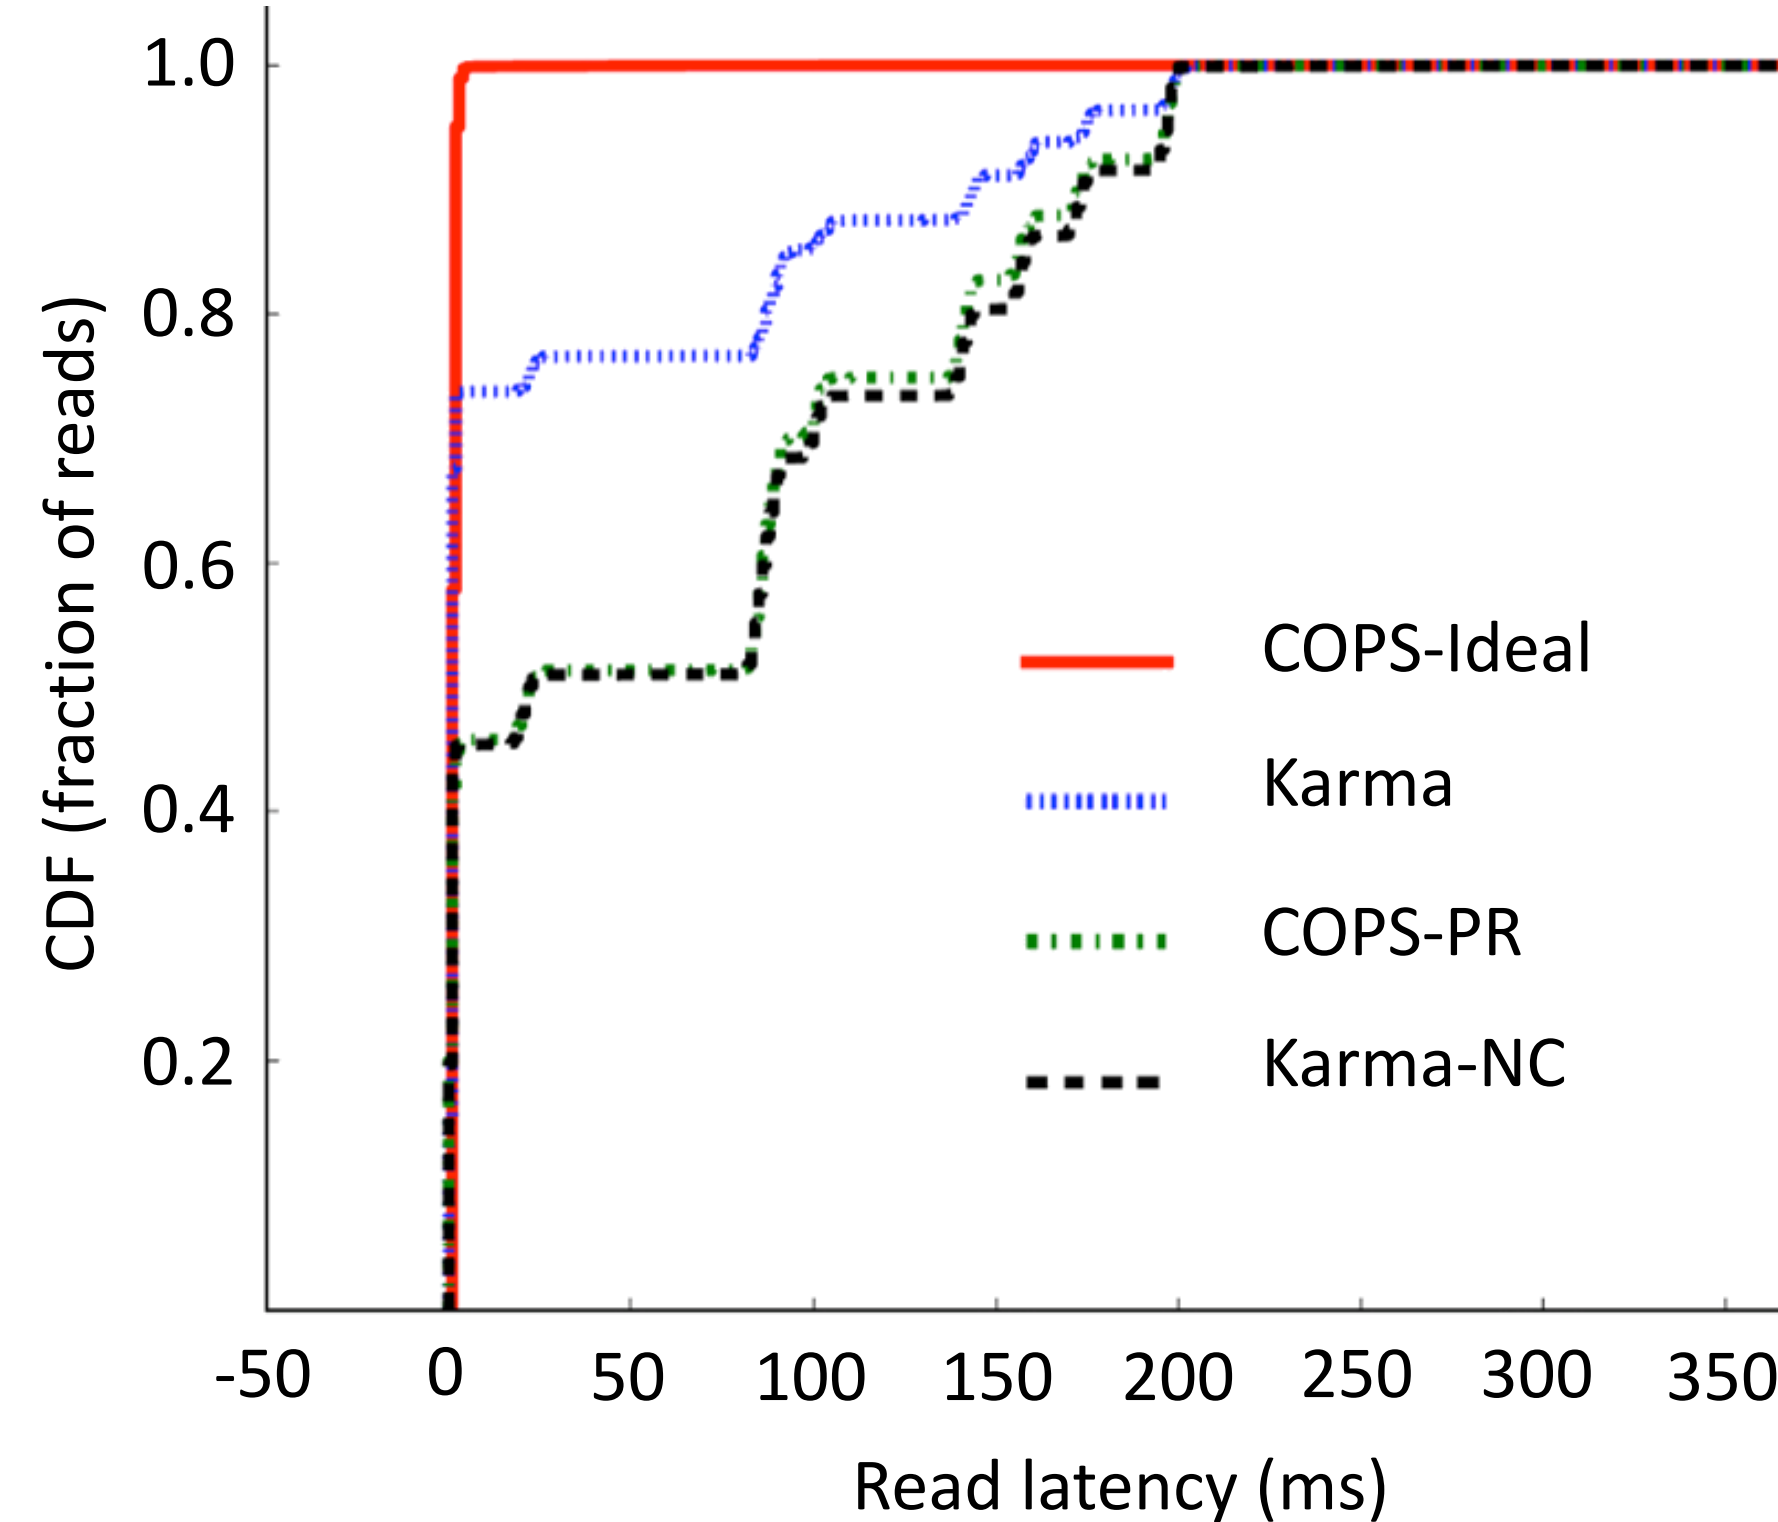
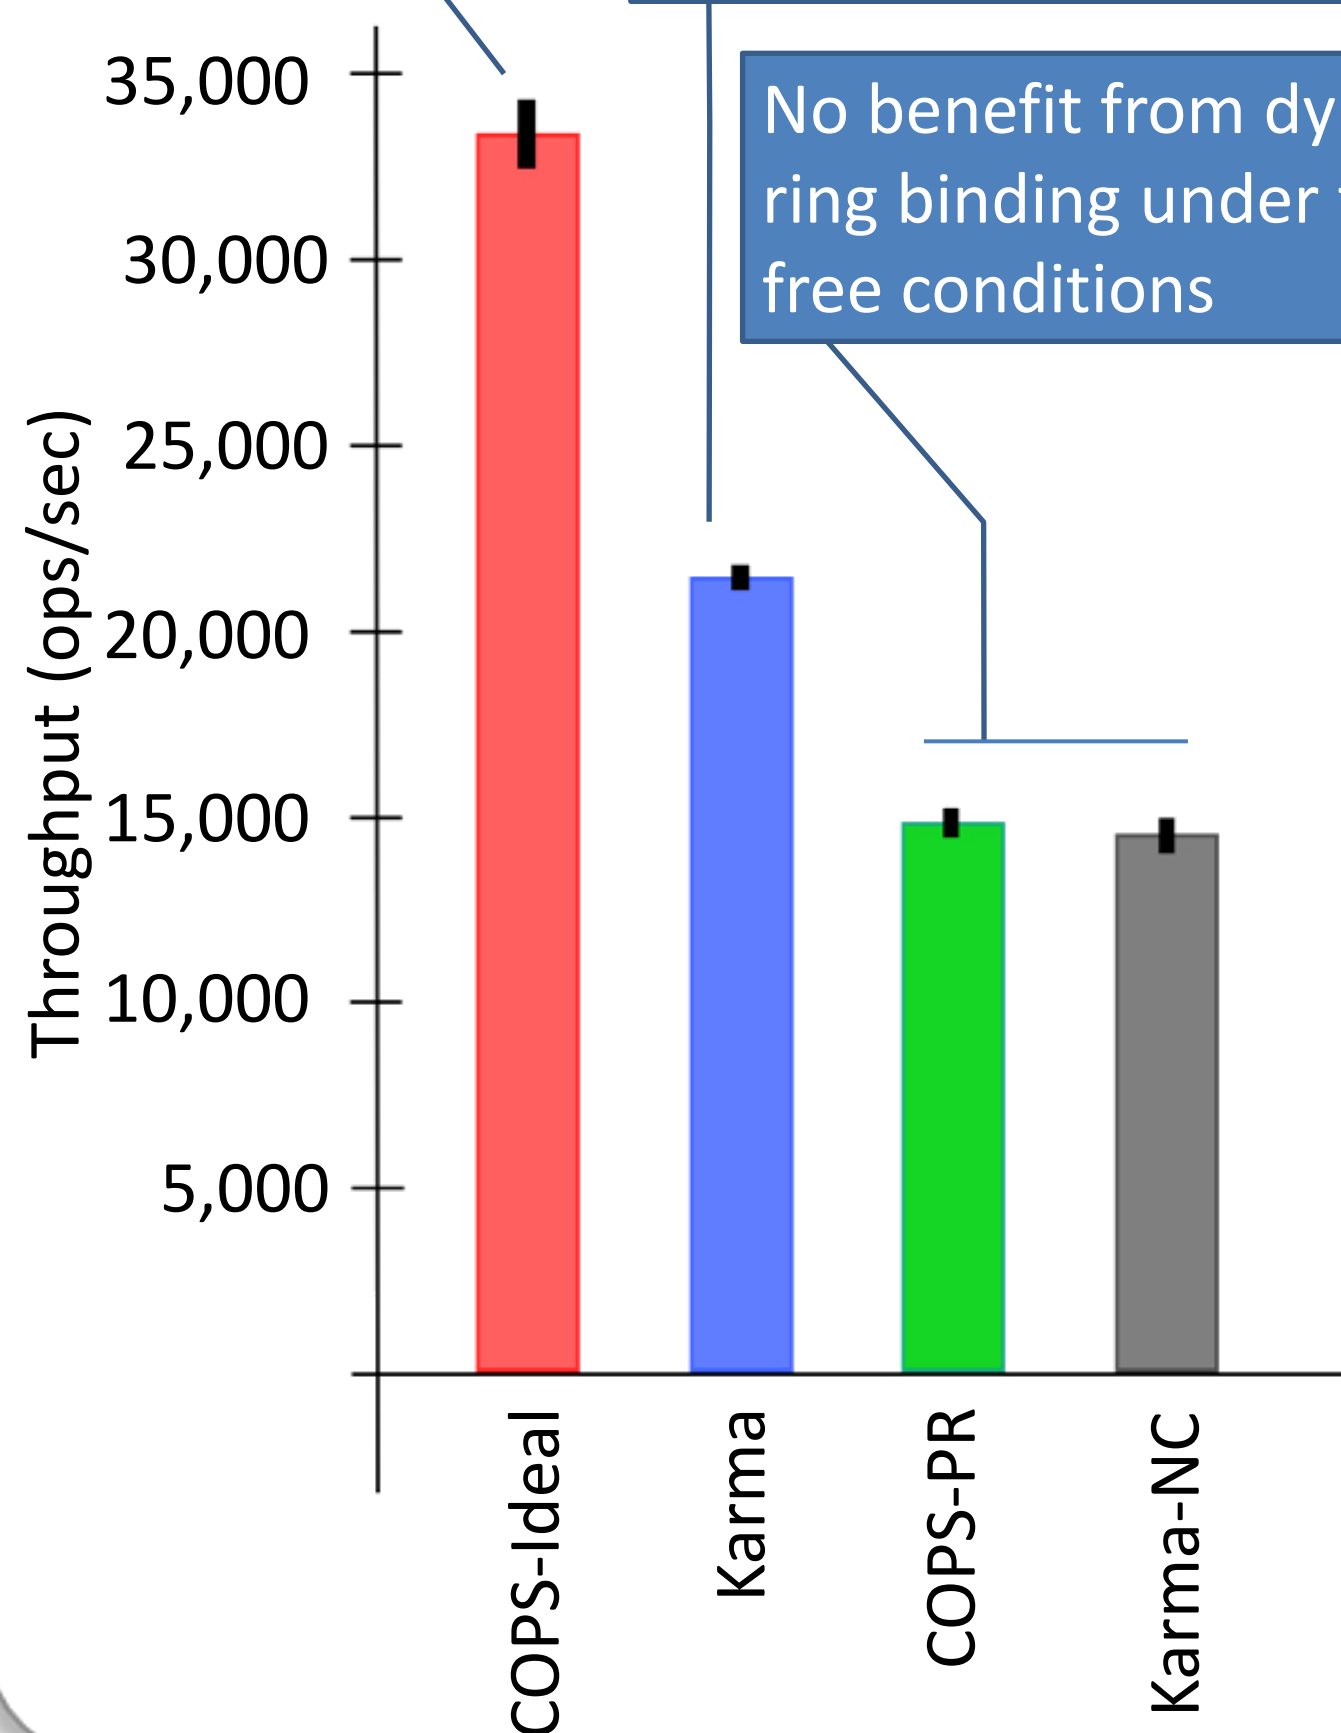
Four Schemes:

	COPS-Ideal	Karma	COPS-PR	Karma-NC
Replication	Full	Partial	Partial	Partial
Ring binding	Static	Dynamic	Static	Dynamic
Write buffers	-	✓	✓	✓
Caching	-	✓	✗	✗

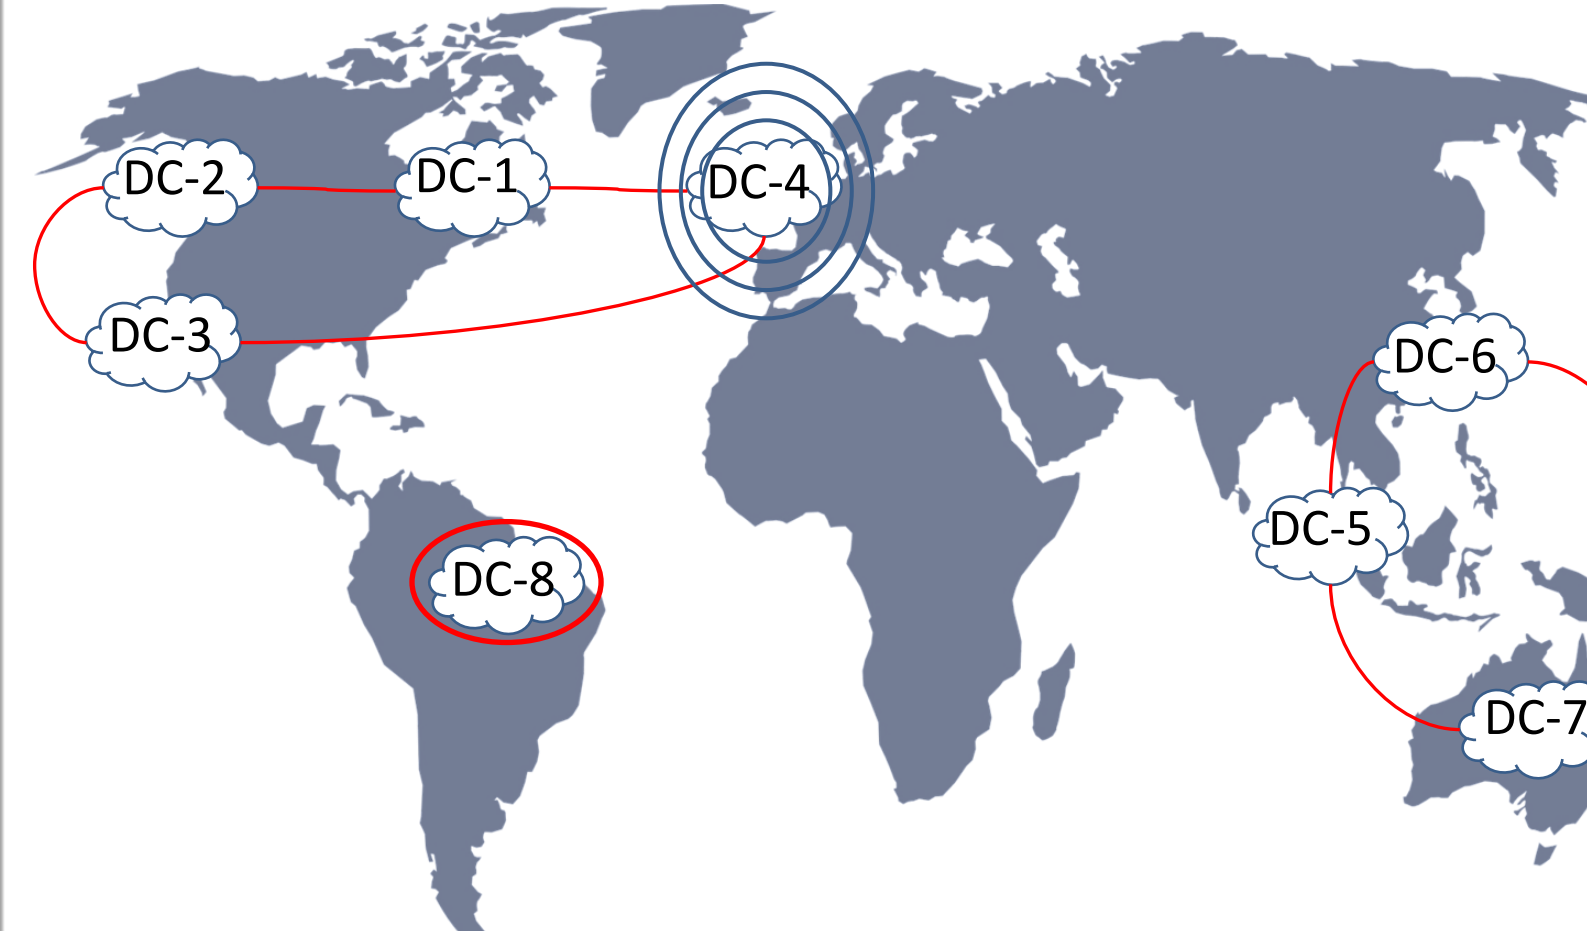
COPS-Ideal looks better, incurs high (impractical) cost of full replication
Karma has $\approx 30\%$ cache misses
Karma's advantage diminished here since writes are only 5%

Karma achieves 43% higher throughput than COPS-PR

No benefit from dynamic ring binding under fault free conditions



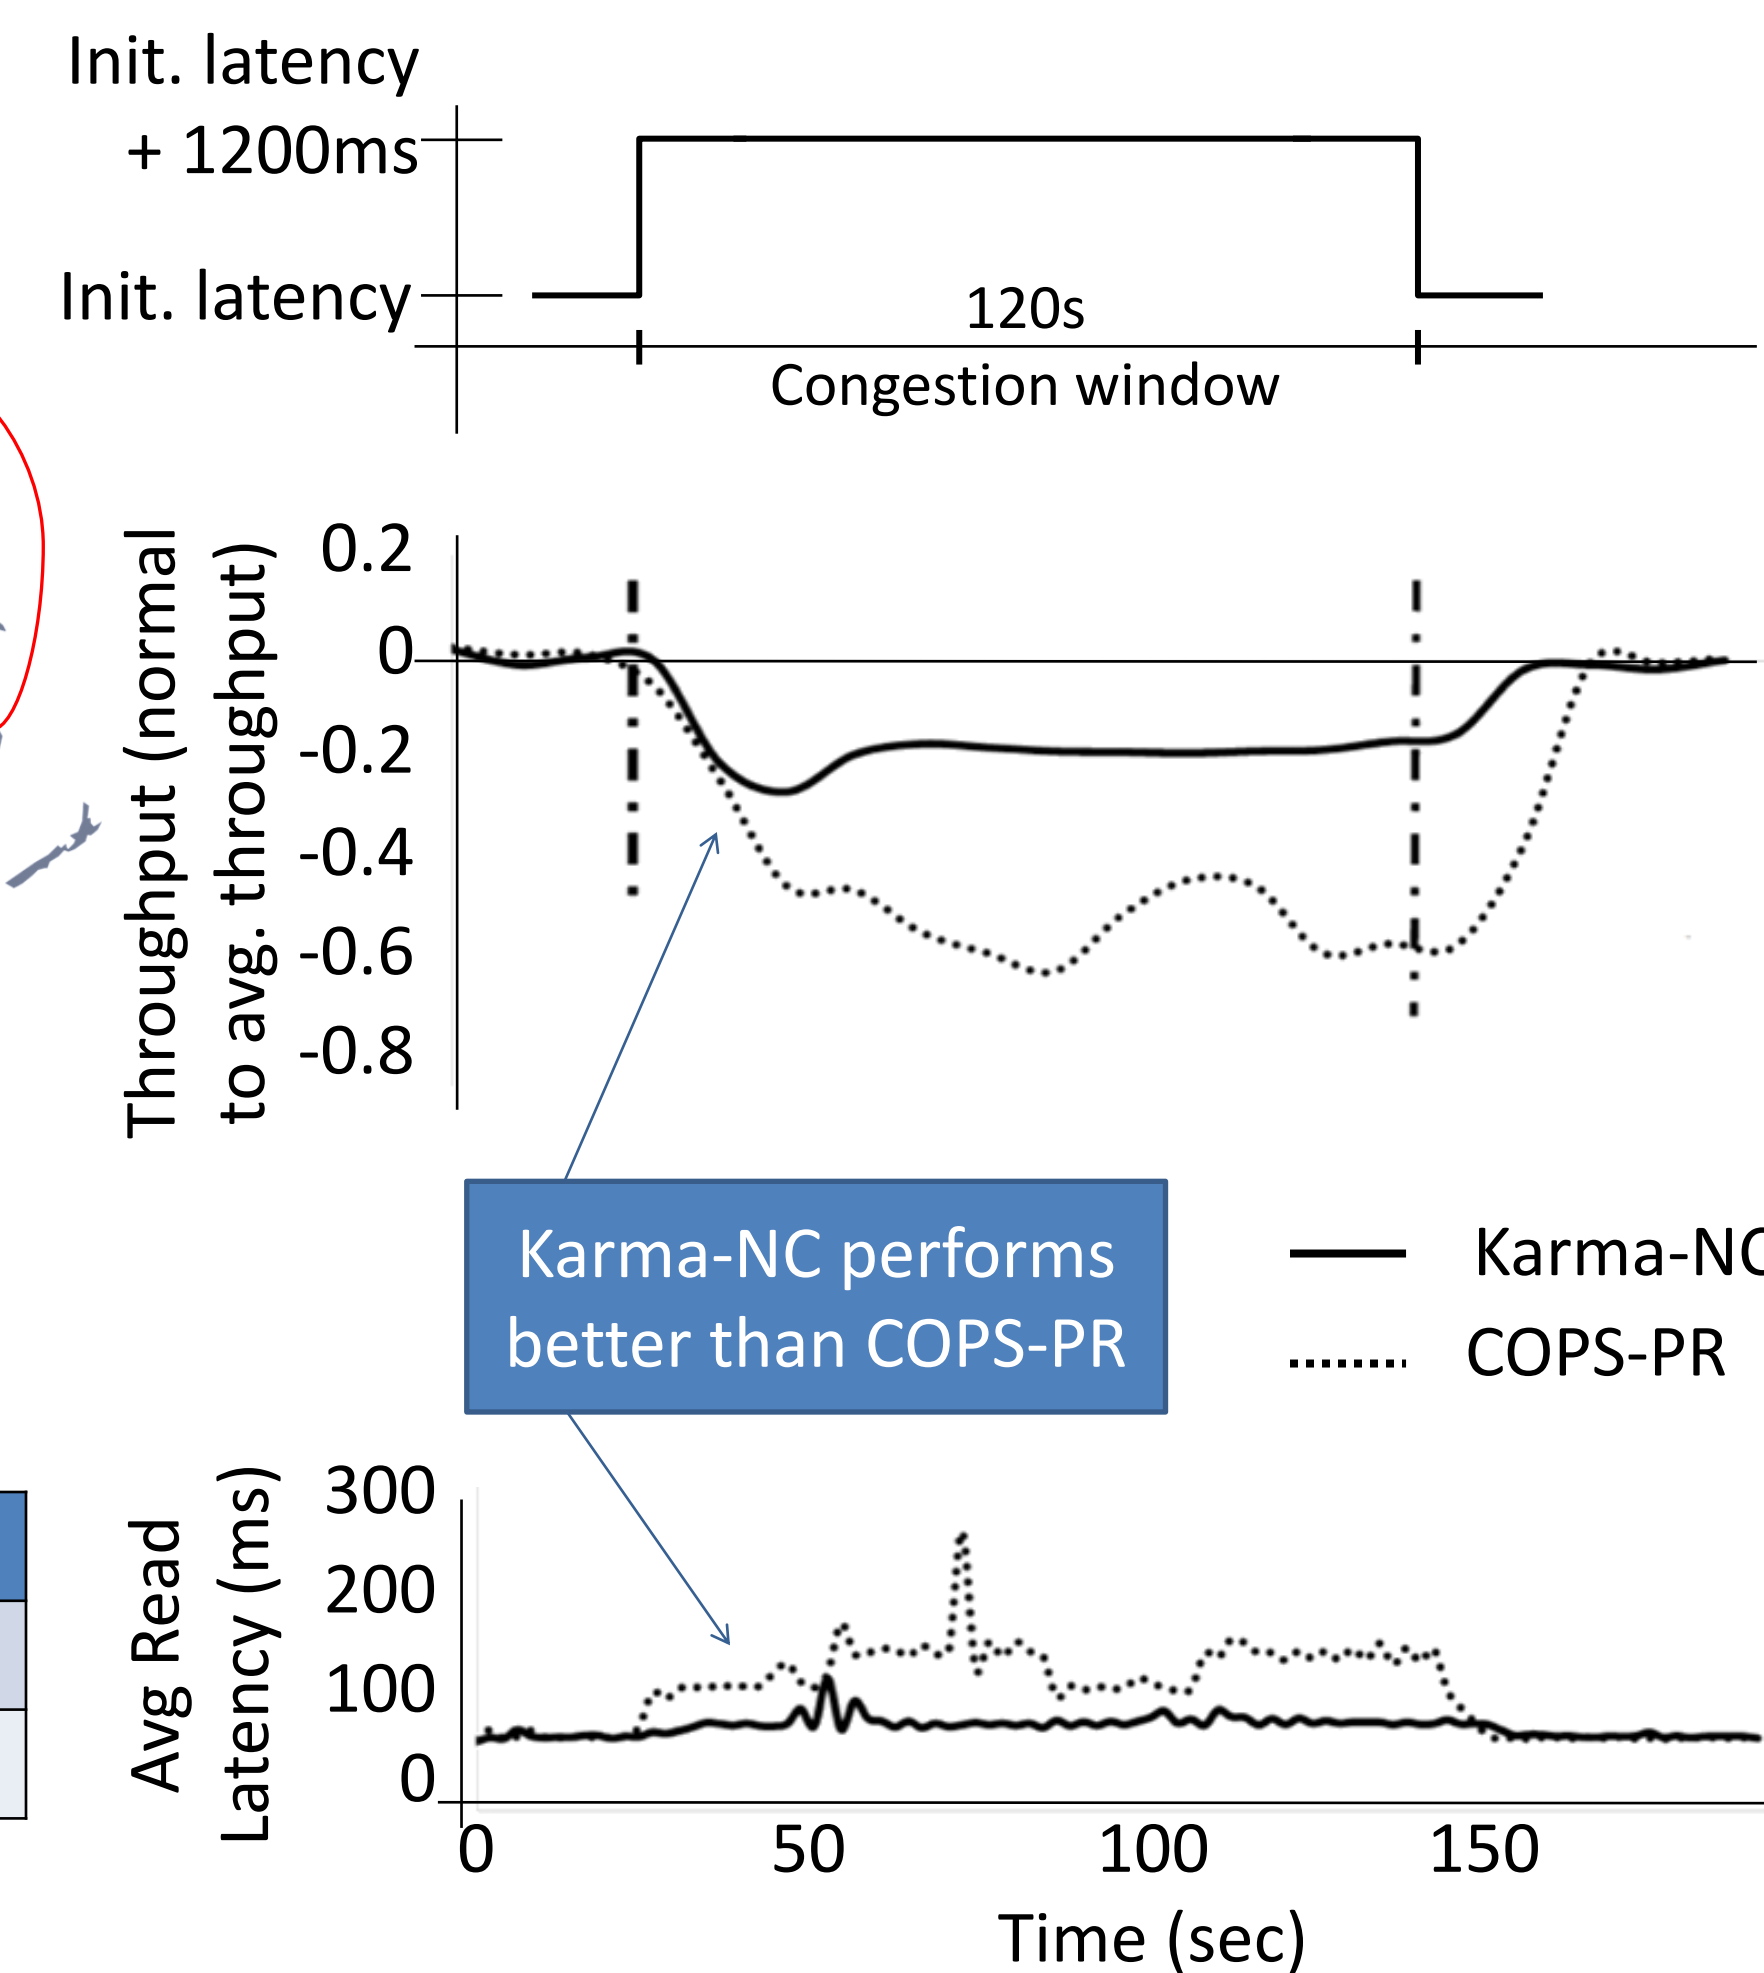
Performance with Faults



Importance of Dynamic Ring Binding

- Induce congestion in Europe DC
- All traffic (in and out) is affected
- Table below shows avg. performance hit



	Karma-NC	COPS-PR
Throughput degradation	$\approx 20\%$	$\approx 50\%$
Latency increase	$\approx 20\%$	$\approx 85\%$



Fault Tolerance Analysis

Failure	Availability	Contrast to Full Repl.	Contrast to COPS-PR	Protection Mechanism
Backend Server	✓	=	=	Chain replication
Cache Server	✓	Not applicable		Stable state
Frontend Server	✓	=	=	Chain replication
Rack	✓	=	=	Chain replication
Single AZ	✓	=	↑	Dynamic binding
Partition	✓	=	↑	Dynamic binding

Cost Savings



Summary

- First causally-consistent cloud storage system with:
 - Partial replication \Rightarrow Practical, cost effective
 - Dynamic ring switching \Rightarrow Stronger availability guarantees
- 43% throughput improvement iso-cost
- Significant reduction in operational and capital expenditures